cLACIr: Cognitive Load and Canine Intervention Recognition Dataset

Walker S. Arce

University of Nebraska-Lincoln, Department of Electrical and Computer Engineering, Lincoln, 68588, United States

E-mail: wsarcera@gmail.com

Jeffrey R. Stevens

University of Nebraska-Lincoln, Department of Psychology, Lincoln, 68588, United States

E-mail: jeffrey.r.stevens@gmail.com

September 2025

Abstract. Emotionally intelligent machines require the ability to differentiate emotional responses from humans from any background. Naturally, each human's emotional responses, whether it's physiological, facial, or any other modality, can be very unique to the individual. While these specific responses are unique though, there are underlying trends and patterns that can be understood when using larger sample sizes. The Cognitive Load and Canine Intervention Recognition (CLACIR) dataset consists of 140 subjects whose physiological responses to cognitive load and intervention are recorded using an Empatica E4. The dataset is split into two groups of subjects which are roughly equivalent in size and separated temporally. Physiological data consists of 3-axis accelerometry, electrodermal activity, and photoplethysmography recorded from the left wrist. This dataset is useful as a baseline for assessing machine learning models as well as for pretraining or mixing with other datasets.

Introduction

Human emotion and affect are physiological and behavioral reactions to changes in our environment [cite]. Some affective states may be easy to label (e.g., anger) but others may be more difficult to pin down (e.g., stress, anxiety). What's more, individuals vary greatly in how they experience affective states, with some experiencing intense emotions and others experiencing more muted emotions. The field of affective computing develops machines and algorithms that can detect human affective states using measures of physiological responses such as heart rate, breathing rate, skin conductance, and muscle activation [**picard2000affective**]. These physiological data are input into algorithms that attempt to classify the affective state experienced when the data were collected.

[Need to give short explanation for why this is important. What does affective computing achieve?]

~~When applying machine learning to a vision-based task, ImageNet—a common dataset for pretraining vision models—contains 14 million images across twenty-thousand categories [**deng2009imagenet**]. The concept of pretraining utilizes the idea that there are underlying structures or patterns that are transferable to most, if not all, classification tasks. In many cases, this is accomplished by training a machine learning algorithm on a very large structured, or unstructured, dataset so that the model can learn useful, generalizable features.~~

Accurately classifying human affective state requires large data sets due to the variability in emotions experienced within and between people. But the number of large-scale datasets for physiological signals is limited. Collecting physiological data can be difficult because participants must wear recording equipment either during their daily life or in a laboratory study [**gjoreski2020datasets**, **spathis2020learning**, **schmidt2018introducing**, **goodwin2019predicting**, **thayer2022effects**, **hosseini2022multimodal**, **poh2010wearable**, **ekiz2020can**]. The human body generates physiological signals through chemical and electrical interactions with various organs, which are largely involuntary and not easily controlled by an individual [**goodwin2019predicting**]. Consequently, targeted emotional states must either be induced in participants within a laboratory setting or annotated through surveys from naturalistic settings. Certain cognitive tasks can be used to reliably induce physiological states. For instance, researchers can put participants into a state of *cognitive load*—where cognitive limits of processing information are strained—by engaging long-term memory, working memory, and attentional control [cite] [~~**thayer2022effects**~~]. Physiological measures show that cognitive load can cause stress-like responses in the autonomic nervous system, including changes in heart rate, heart rate variability, and galvanic skin response (Portnova et al. 2023; `https://doi.org/10.1007/s11055-023-01394-9`). However, previous work has found ways to discriminate cognitive load from stress. Notably, Setz et al. 2010 found ... [Need to add in what we already know about detecting cognitive load from Setz et al 2010 `https://doi.org/10.1109/TITB.2009.2036164` and Markova et al. 2019 `https://doi.org/10.1109/BIA48344.2019.8967457`]. ~~Currently, this limitation necessitates feature engineering, where humans decide which features to extract from a dataset to feed into a machine learning model [**picard2000affective**, **spathis2020learning**, **schmidt2018introducing**, **goodwin2019predicting**, **t**~~

[Sorry, but this part below feels like justification for methods rather than an introduction to the study.]

~~Due to a lack of larger datasets, many studies utilize the leave one subject out (LOSO) approach, which excludes one subject as the testing set and uses the remaining as the training set [**schmidt2018introducing**, **farrahi2019evaluating**]. This has preference because it maximizes the amount of data used in training while also implementing an identity split, where a separate person, whom the model has~~

~~never seen, is tested. It is possible that the excluded subject is the subject that the trained models perform the best on, which would be an unrealistic estimation of both the amount of variance between subjects and the generalizability of selected features [**mac2002problem**, **li2006reduction**]. One method to remedy this is to balance the representation of each subject and each class and to perform multiple splits of the dataset based on these representations, rather than just each class [**li2006reduction**]. This method is called a stratified, group, k-fold split, which balances the class representation and the subject representation across k-folds [**pedregosa2011scikit**].~~

~~Additionally, collecting physiological datasets can be difficult because subjects must either wear equipment during their daily life or participate for a short period using laboratory equipment [**gjoreski2020datasets**, **spathis2020learning**, **schmidt2018introducing**, **good~~
~~The time-series physiological signals being measured are generated by the human body through chemical and electrical interactions with various organs, which are largely involuntary and not easily controlled by an individual [**goodwin2019predicting**]. Consequently, targeted emotional states must either be generated in subjects within a laboratory setting or annotated through surveys from naturalistic settings by the subject. Cognitive tasks can also be used to reliably generate physiological signal changes in subjects, using tools such as the Deese-Roedinger-McDermott task to test long-term memory, the Necker cube pattern control to test attention control, the backwards digit span to test working memory, and the n-back to test working memory [**thayer2022effects**]. The labor required for these and other methods contributes to the low subject counts as well as to a low number of emotive states within the recorded data.~~

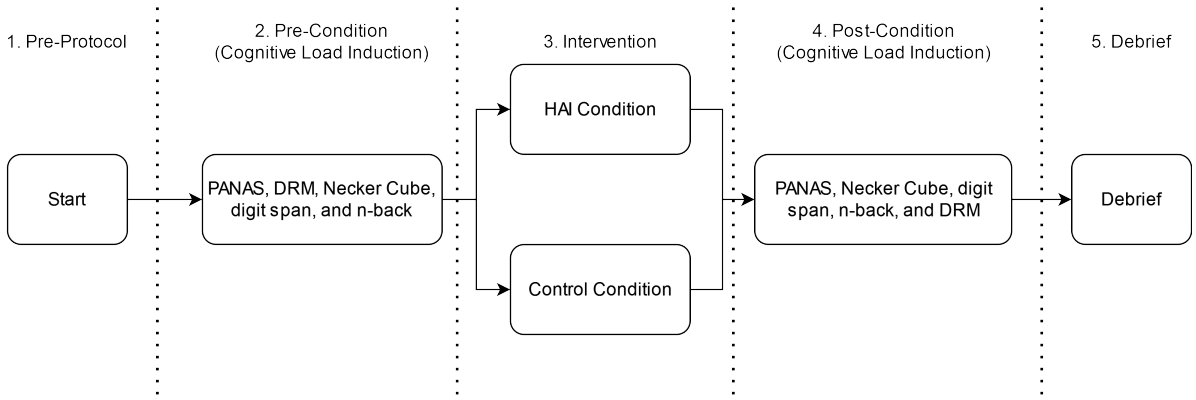

Figure 1. The experimental protocol to induce cognitive load states and provide intervention from the same using a control and human animal interaction (HAI) conditions.

Here we introduce the Cognitive Load and Canine Intervention Recognition (cLACIr) dataset, which provides multimodal physiological data from a variety of mental states including active cognitive load, control intervention, active intervention, and post-intervention cognitive load. The dataset comes from a study of how human-animal interactions influence mood, stress, anxiety, and cognition [**thayer2022effects**].

The study tested 140 participants—all of whom were undergraduates in the University of Nebraska-Lincoln Department of Psychology subject pool—in three experimental phases: pre-condition, intervention, and post-condition (Figure 1). The pre-condition phases tested participant long-term memory (Deese-Roedinger-McDermott task), attentional control (Necker cube pattern control task), and working memory (backwards digit span task and n-back task). After completing these cognitive tasks, participants entered the intervention condition in which they either interacted with a dog (active intervention) or completed a repetitive task (control intervention) for three minutes. After the intervention, participants completed the same cognitive tasks in a differing order.

The **thayer2022effects**aper analyzed the cognitive and mood data but not the physiological data collected during the study. We collected physiological data via an Empatica E4 wristband [**poh2010wearable**], including blood volume pulse, electrodermal activity, temperature, and movement. This dataset includes 95.8 hours of these four recordings across 140 participants.

Due to the size of this dataset, large identity splits can be achieved, providing a fair method of comparing different machine learning architectures. With cLACIr as a foundation, clinical applications can use machine learning for pilot studies in addition to traditional statistics[**phippsassessing**], more robust machine learning models can be incorporated into digital domains like virtual reality [**arce2021biosensor**], or can improve aggression detection approaches [**goodwin2019predicting**, **arce2023detecting**].

Methods

Data Collection

This dataset was collected as a series of two experiments to investigate the effect of human animal interactions (HAIs) on affect and cognitive ability of the participants. The first experiment was conducted from September to November 2018 and recruited 73 subjects, 13 of which identified as male (17.8%) and 60 identified as female (82.2%), with a mean age of 19.2 (SD = 1.4) years old. The second experiment was conducted from November 2018 to April 2019 and recruited 83 participants, 17 of which identified as male (20.5%) and 66 identified as female (79.5%), with a mean age of 19.9 (SD = 1.8) years old. All participants were recruited from the University of Nebraska-Lincoln's psychology subject pool and did not have a physical or emotional aversion to dogs.

2 is a schematic representation of the hierarchical structure of the dataset.

To collect this data, an experimental protocol was designed that would induce a state of cognitive load on the participant, then apply an intervention condition, and then apply another cognitive load condition. The cognitive load induction has four separate cognitive tasks: Deese-Roedinger-McDermott (DRM) to test long term memory[**roediger1995creating**, **mcevoy1999connection**], the

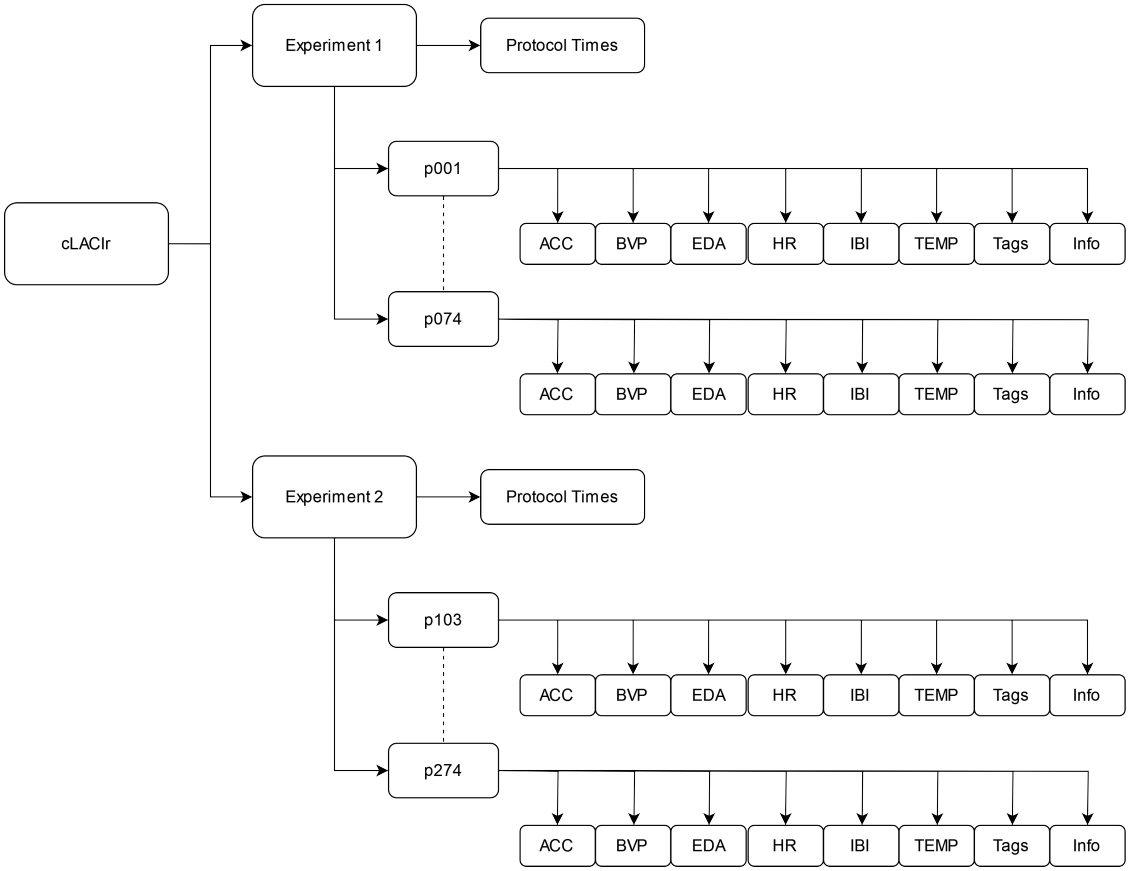

Figure 2. The hierarchical structure of the dataset.

Necker cube pattern control to test attentional control[**orbach1963reversibility**, **cimprich1993development**, **sahlin2016influence**], the backwards digit span to test working memory[**berman2008cognitive**], and the n-back to test working memory[**rich2007effects**, **cohen1994activation**]. These assessments were presented and responses collected using PsychoPy version 1.90.2[**peirce2019psychopy2**] running on a computer's 16-inch monitor with only the participant and researcher present.

After the cognitive load state is induced in the participant, they move on to the intervention condition in a separate room. Each participant is pseudo randomly assigned to either a human animal interaction (HAI) condition or a control condition. For the HAI condition, the participant interacts with the second author's dog, who is a 65-pound, neutered, male Catahoula leopard mix that was Canine Good Citizen certified prior to the study. For the control condition, the participant is alone in the room and provided a sheet of paper with a full page of Latin text and instructed to circle every "e" and "f" for the "next few minutes". The participant was informed that the task would not be graded and there is no penalty for wrong answers. The intervention condition was performed for three minutes in both cases, based on the results from [**handlin2011short**]. After the intervention condition, the participant is moved on to a second cognitive load induction activity.

To measure the participant's affect, the Positive and Negative Affect Schedule (PANAS)[**watson1988development**] was administered before each cognitive load induction. The second experiment also added additional measures of anxiety, including present and general feeling of anxiety with the State and Trait Anxiety Inventory (STAI) from [**spielberger1999measuring**] and their present feelings of anxiety using the Anxiety Visual Analogue Scale (AVAS) from [**cella1986reliability**]. Due to the new measures, the cognitive load induction tasks were modified to perform the AVAS first, then STAI, and then PANAS.

To measure the physiological responses of the participants as they go through the study protocol, which were not analyzed in previous studies due to low sampling rate of the heart rate signal [**shaffer2017overview**, **thayer2022effects**, **Stevens2020EffectsOH**], an Empatica E4 biosensor is affixed to their left wrist. This wrist-worn biosensor has been used to generate data in many studies [**goodwin2019predicting**, **thayer2022effects**, **schmidt2018introducing**] and is regulatory compliant and validated [**poh2010wearable**], as well as being the top rated device for patient populations served by the behavioral analysis clinics [**taj2018review**]. This biosensor has four sensors: 64 Hz blood volume pulse (BVP) using photoplethysmography (PPG), 4 Hz electrodermal activity (EDA), 4 Hz infrared thermopile, and a 32 Hz 3-axis accelerometer measuring accelerations in its inertial frame of reference [**goodwin2019predicting**]. When transitioning between states of the protocol (start -> cognitive load induction -> intervention -> cognitive load induction -> debrief) the button on the biosensor is pressed, placing a marker in the datastream. This marker is used to label each data point with its associated protocol stage during post-processing.

The E4 has multiple methods to compensate for common issues with in-situ physiological signal recording. Measurement error in the PPG signal due to skin color and changing light intensity is compensated dynamically by the device [**goodwin2019predicting**]. Motion artifacts are removed from the captured signals and heart rate variability artifacts are removed as well [**goodwin2019predicting**]. Having this preprocessing be performed on the data before it is reported to the collection software reduces the burden on preprocessing for this experiment. Due to this, minimal preprocessing is needed before training computational intelligence models.

Each physiological signal from the Empatica E4 encodes one of two response modalities from the subject. The first is the physical movement of the subject encoded by the accelerometer and the second is the autonomic nervous system responses to the environment. By looking at the movement of the subject's hand, specific movements that occur during emotional states can be assessed, which is the basis of physical activity recognition [**ravi2005activity**]. The separability of movements from different stimuli will likely be highly variable between subjects and may be highly correlated to the level of reaction [**thang2012gait**].

Alternatively, the autonomic nervous system responses are highly correlated to internal emotional state [**kreibig2010autonomic**] and is the basis of emotion

recognition in machine learning systems [**healey2005detecting**]. The most common physiological signal used for this task is the EDA signal, which is collected by measuring the resistance between two contacts pressed against the subject's skin, which is the same principle as a lie detector [**society2012publication**, **arce2021biosensor**]. It can be decomposed into three constituent signals: the skin conductance response (SCR), the skin conductance level (SCL), and noise [**greco2015cvxeda**]. The noise will be ignored. The SCR is the fast-changing response to events, constituting a phasic component of the EDA signal, and can be used to approximate sympathetic nervous system activity [**society2012publication**, **benedek2010continuous**]. The SCL is a slower changing, tonic component of the EDA signal that acts as the baseline conductance of the skin. These two components indicate how strongly an emotion was felt and can also be used to differentiate between emotional states, for instance, the SCL and SCR will go up when someone is feeling anxiety and only the SCL will go up if someone is feeling happiness [**kreibig2010autonomic**].

Peripheral skin temperature is a common measure of stress and can be used to distinguish between emotional states [**collet1997autonomic**, **salazar2015mental**, **karthikeyan2012descriptive**]. Generally, the theory of blood flow constriction due to stress causing a decrease in skin temperature and blood flow dilation due to happiness causing an increase in skin temperature is used [**karthikeyan2012descriptive**]. But there are conflicting reports of skin temperature response to emotional states, where few report the temperature of the wrist or hand and instead focus on either monitoring facial temperature with a thermal camera [**salazar2015mental**] or will place their temperature probe in an area that is free from thermal environmental disturbances, such as under the armpit [**karthikeyan2012descriptive**]. For studies that place the temperature probe on the hand [**collet1997autonomic**], a decrease in skin temperature is observed during happy states and an increase in skin temperature is observed for anger. Considering the placement of the Empatica E4, the observations from the temperature probe on the hand will be considered.

Finally, PPG is an indirect method of measuring the cardiac cycle and can be compared with the electrocardiogram (ECG), which directly measures the electrical activity of the heart [**van2019heartpy**]. PPG is performed by measuring the change in skin coloration as blood travels through the capillaries of the skin, which is directly correlated to the blood being pumped by the heart. From PPG, heart rate in beats per min (BPM) can be derived. In the case of someone feeling anxiety and happiness, their heart rate will go up [**kreibig2010autonomic**].

The results of these methods and the statistical data analysis of the surveys are analyzed in [**thayer2022effects**] for the first experiment and in [**Stevens2020EffectsOH**] for both. The physiological data has been analyzed and used for machine learning in the first author's master's thesis [**arce2023unobtrusive**].

Data Preprocessing

To process the dataset, handcrafted features will be extracted from overlapping 60-second windows of data. These features consist of statistical information generated using the FLIRT toolbox [**foll2021flirt**]. The EDA signal is decomposed into its constituent signals using cvxEDA [**greco2015cvxeda**], where the noise is discarded and the SCL and SCR are processed using FLIRT, in addition to the original signal. The other signals (PPG, 3-axis accelerometry, and peripheral skin temperature) are processed in their original form using FLIRT. The calculated statistics for each signal window are the mean, standard deviation, minimum, maximum, dynamic range, sum, energy, skewness, kurtosis, peak count, RMS, line integral, number of samples above and below mean, number of sign changes, 25%-75% interquartile range, 5%-95% interquartile range, 5-th percentile, and 95-th percentile. Additionally, the entropy, perumutation entropy, and singular value decomposition entropy are calculated. All of these metrics are handled by FLIRT[**foll2021flirt**]. Each signal is broken into 5-second windows, which is a common approach [**schmidt2018introducing**, **phippsassessing**], and those windows have their statistical features calculated. Each participant's data will be tagged with their participant number, allowing the identity of the data to be tracked during training.

The labels are assigned based on the protocol steps for the first question, where the control intervention is considered neutral (label 0), cognitive load induction is considered stress (label 1), and HAI intervention is considered amusement (label 2), to align with [**schmidt2018introducing**]'s labeling. For the binary task, the cognitive load class and HAI intervention are merged. To assess how induced cognitive load can be reduced through different intervention conditions, a second question that will be posed is whether the machine learning models can differentiate between post-condition that are preceded by the HAI intervention or control intervention. The idea behind this question is that the level of cognitive load being experienced by the participant (and thereby recorded by the biosensor) is reduced in comparison to the control condition. Both of these questions will be repeated for the condition where the accelerometry features are removed. By including these features, we hypothesize that it is likely to be a naive discriminator for the models, where HAI can be detected by wrist movements caused by petting the dog.

Machine Learning Approaches

The dataset will be benchmarked using a standard selection of common machine learning models. These models are Linear Discriminant Analysis (LDA), Decision Tree Classifier (DTC), Random Forest Classifier (RFC), and the K-Nearest Neighbors algorithm (KNP), all of which are implemented in SciKit-Learn [**pedregosa2011scikit**]. An LDA model uses a linear combination of features to separate two or more classes, where predictions are made using Bayes' rule where the selected class maximizes the posterior probability [**hastie2009elements**]. These models are also closely related to analysis of variance (ANOVA), but the LDA uses continuous independent variables and categorical dependent variables [**wetcher2011analyzing**]. The DTC model uses

simple decision rules, such as feature a is less than feature b infers class a, which makes predictions very explainable [**breiman1984nonlinear**]. The RFC is a meta-estimator that uses an ensemble of DTCs to predict classes based on the input data by taking a majority vote from the DTC ensemble [**breiman1999random**]. Finally, the MLP is a neural network that multiplies the input data by an input weight matrix, which is passed through an activation function and then finally multiplied by an output weight matrix, which produces several outputs proportional to the data classes [**hinton1990connectionist**]. These models are trained by adjusting the weight matrices using various backpropagation algorithms, such as the Adam optimizer that moves the operating point of the model using first- and second-order moments [**kingma2014adam**].

To train the machine learning models, a stratified group 10-fold approach will be used [**pedregosa2011scikit**]. Using the identity labels, dataset features, and dataset labels, 10 splits of the data will be generated, where each split contains a different group of identities and the number of samples from each class is as balanced as possible. By doing this, not only will the test set contain unique identities from the train set, a requirement for generalization performance, but the choice of identity for testing will also be evaluated. Doing this, a mean and standard deviation for each evaluation metric will be generated. If there are ideal and non-ideal choices for test participant, then this will be observed as a high standard deviation in the final metrics. Finally, the PANAS results are used to filter out subjects in the dataset by thresholding the difference between pre- and post-condition PANAS scores. By varying this from -1 to 1, we can filter out subjects that had an increasingly positive reaction to the data collection and see if the physiological data reflects this.

Figures of Merit

While the training of machine learning models can be simplified to a few function calls using highly abstracted software libraries, the methods to compare trained models in a fair way is still an active question. The most common method of comparison is accuracy, or more properly, top-1 accuracy. This metric quantifies how often the highest probability class is the target class. There is also top-k accuracy, which looks at an increasing number of output probabilities and, if the target class is in that list of probabilities, then it counts as correct (Liu, 2014). For binary tasks, there is only top-1 accuracy.

The next most common metrics are recall- and precision-based metrics. Recall as a metric quantifies how often the model correctly chooses a class versus the number of times that it should have chosen that class (Davis, 2006; Liu, 2014), while precision quantifies how often a class is chosen correctly versus all the times that class is chosen (Davis, 2006; Liu, 2014). These two metrics can be combined to create derivative metrics, such as the F1 score, which is the harmonic mean of precision and recall (Liu, 2014) and can give insight into how well a model balances those two metrics. At this

level, these mentioned metrics are common across all machine learning tasks.

Computer vision has utilized many robust metrics that utilize receiver operator characteristic (ROC) curves to quantify the diagnostic ability of a binary classifier at varying threshold levels on the probability outputs (Davis, 2006; Fawcett, 2006; Liu, 2014). For these curves, the true positive rate (TPR) is plotted on the y-axis and the false positive rate (FPR) is plotted on the x-axis for all values of the threshold from 0 to 1. An optimal relationship between TPR and FPR occurs when the TPR is very high and the FPR is very low, meaning that the curve approaches the upper left corner of the plot. The area under the curve (AUC) of the ROC curve is a way of capturing the curve in a single number, where the closer it is to one, the closer the ROC curve is to the upper left of the plot area. Setting thresholds on the FPR can be useful for establishing a confidence level in the performance of a model. If the FPR is set to 1%, then a corresponding TPR of 95% would indicate that for every 1 false positive (FP) prediction, 95 true positive (TP) predictions would be accomplished on average. This is commonly extended to two or three thresholds, namely: 1%, 5%, and 10%. Equally important to mention is the equal error rate (EER), which is the ratio of the FPR to the false negative rate (FNR) by drawing a line from the top left corner of the ROC curve to the bottom right corner and finding the intersecting values (Jothilakshmi, 2016). In this case, a lower value indicates that the classifier is performing better, and this metric is important when the risk of false positive or false negative identification is equally harmful.

While ROC metrics such as AUROC are used in affective computing research, it is not widespread and isn't accompanied by other ROC metrics, such as TPR @ x% FPR, which would provide heightened insight into the confidence of a classifier. Other metrics are being discussed for affective computing, namely the area under the precision recall curve (AUPRC) and Cohen's kappa coefficient [**d2018affective**]. AUPRC is constructed similarly to the AUROC, but precision is plotted on the y-axis and recall is plotted on the x-axis for all values of the threshold, when both metrics are high then the curve will increase towards the upper right corner of the plot area. The larger the area under the curve, the greater the precision and recall at all thresholds. This metric is also called average precision (AP) and in multi-class problems, the mean AP (mAP) is the average AUPRC across all classes. Cohen's kappa coefficient quantifies the inter-rater reliability, which is the degree of agreement between independent observers of the same phenomena, which can be considered a more robust accuracy measure because it takes the probability of random chance agreements into account (Cohen, 1960). The concern of not using descriptive metrics is related to the way that affective computing tasks are constructed. In most cases, the classification task of emotion vs. neutral has a data distribution that leans towards neutral (Schmidt et al. 2018; D'Mello, 2018). Consequently, a model trained on this imbalanced data, i.e., 80% neutral / 20% emotion, then predicting zero (neutral) for every sample will generate a top-1 accuracy score of 80% automatically (D'Mello, 2018). Increasing in popularity, balanced accuracy seeks to improve top-1 accuracy by using the TPR plus the true negative rate (TNR) divided

by two, which accounts for some of the class imbalance problems.

Results

HAI vs. Control Intervention Classification

The first classification test of the dataset was differentiating between the HAI condition and the control condition for the intervention stage of the data collection. The results of this are shown in 3. The baseline value for balanced accuracy is shown by the black line, which is the ratio of one class over another.

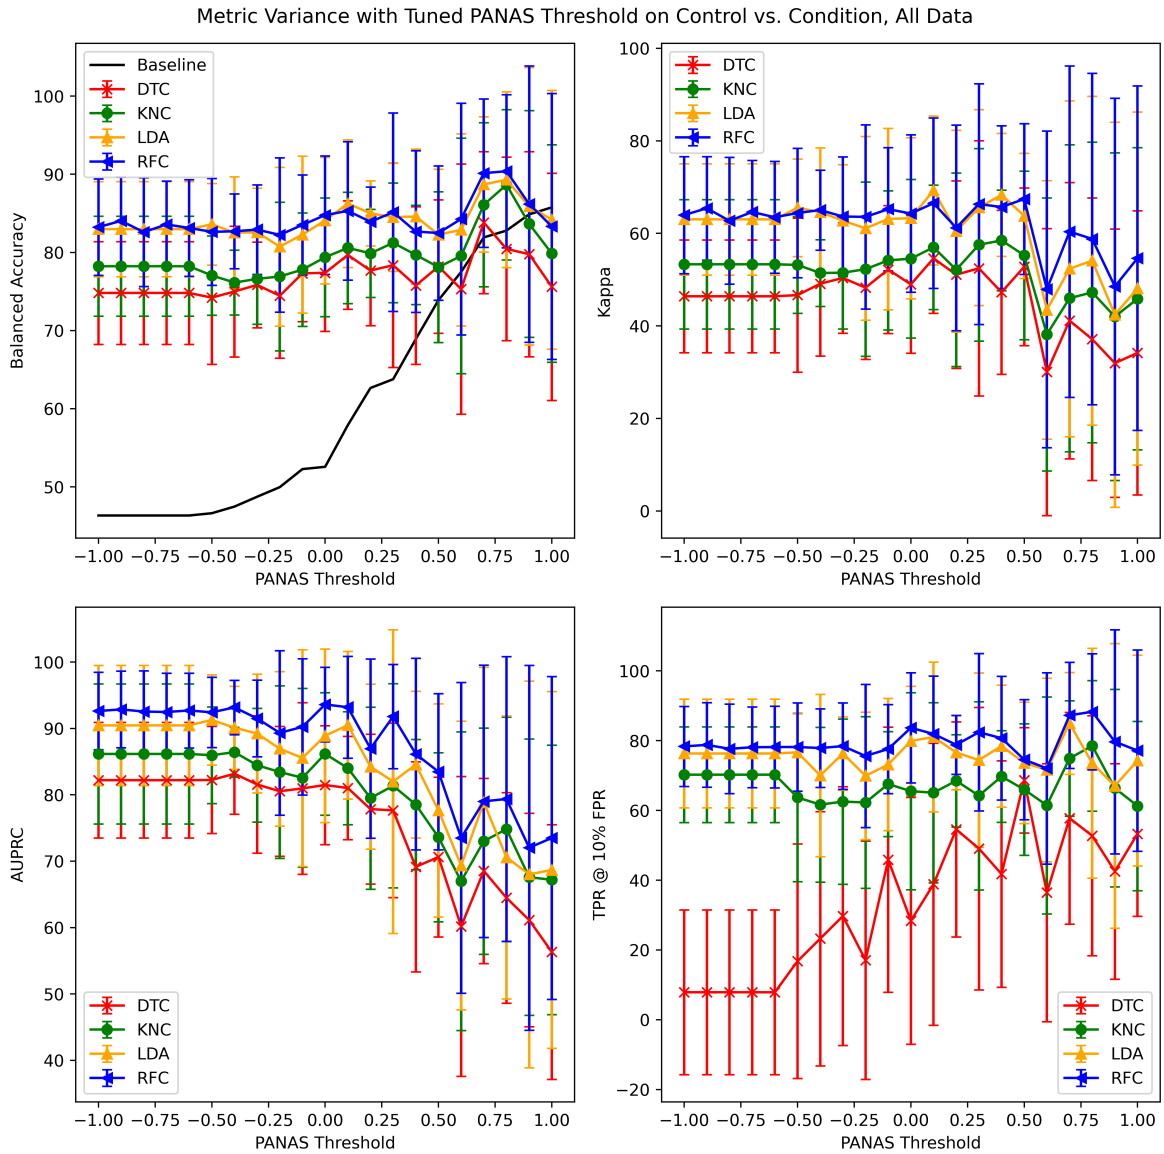

Figure 3. The effect of using various thresholds on the PANAS score to exclude subjects from dataset and training to differentiate between intervention conditions using data from intervention stage.

With the full dataset (PANAS threshold of -1), there are clear improvements that can be achieved by selecting different models. The RFC and LDA perform the best of the four, with balanced accuracy staying above 80% during the ablation study. There is even a slight bump to the balanced accuracy results when the PANAS threshold increases to around 0.5, indicating that removing subjects who had a more negative reaction to the study improves the ability to differentiate between classes. This is contrasted by the other metrics, where AUPRC continues to drop as the PANAS threshold increases. The Kappa test shows some loss in agreement between raters, but the TPR @ 10% FPR shows a similar trend to the balanced accuracy, where it stays solid until the PANAS threshold is sufficiently high and it gets a boost. The models all show similar comparisons to one another across metrics (no model outperforms another in one metric that it underperforms in another), which is not entirely unexpected. Interestingly, the DTC model has a horrible TPR @ 10% FPR, which steadily increases until as the PANAS threshold is tuned.

HAI vs. Control Postcondition Classification

The second classification test of the dataset was differentiating between HAI and control conditions for the postcondition stage of the data collection. The results of this are shown in 4.

This classification task is much more difficult, because subjects are no longer in their intervention stage, so if they are in an excited state, then they are relaxing from it. This is reflected in the data, where the balanced accuracy for this task hovers slightly above the baseline performance. Even as the PANAS threshold is tuned, poor performance continues until around a PANAS threshold of 0.75, where a spike occurs. It doesn't reach the baseline value, but it increases to nearly 70% with all models. The same trends occurred as in 3 for the other metrics.

Discussion

Our dataset performs at the same level as other similar datasets. WESAD, with 15 participants, achieves upper 70% and lower 80% top-1 accuracy results across the board, the DT model performing the worst [**schmidt2018introducing**]. [**gjoreski2020datasets**] had a wider variety of top-1 accuracy values, but achieved a peak value of 82.3% using their Snake dataset and 68.2% using their cognitive load dataset, each with 23 participants. These two examples shows that cLACIr either outperforms or is on par with other similar datasets. Through the sheer number of participants, our models show generalization in our task that is not able to be explored in other datasets due to limited participant count.

It is possible that there is a "cooldown" factor after a participant experiences the HAI condition, such that breaking up the postcondition classification by time bins would improve the accuracy. It is also possible that other machine learning models could

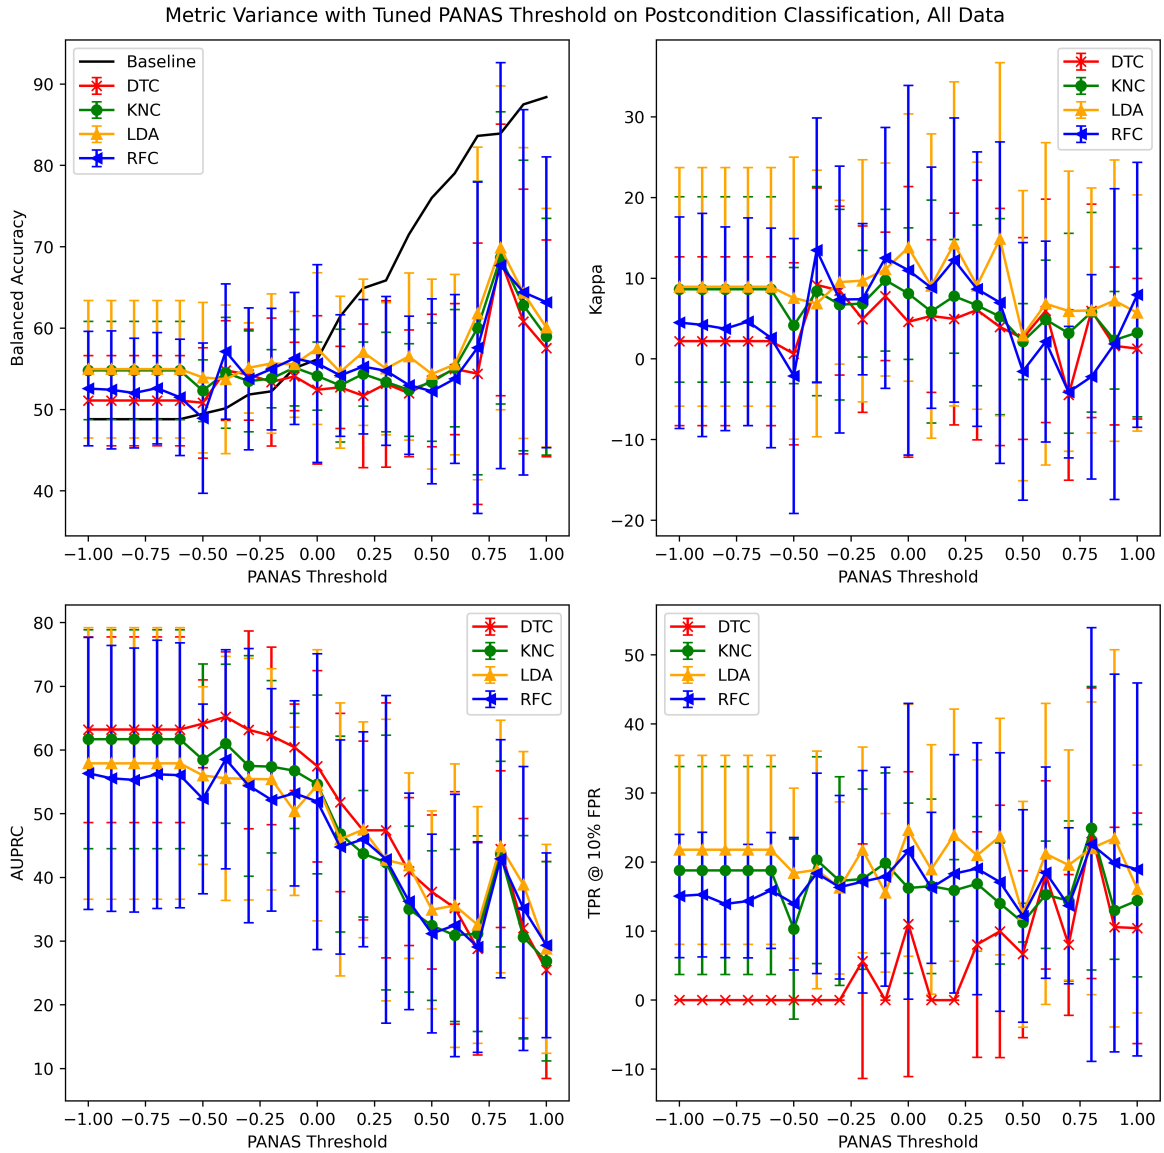

Figure 4. The effect of using various thresholds on the PANAS score to exclude subjects from dataset and training to differentiate between intervention conditions using data from postcondition stage.

perform better, for instance [**gjoreski2020datasets**] uses XGBoost to achieve their top performing model on their Snake dataset. Additionally, there is the question of whether models trained with the cLACIr dataset would generalize to other datasets. Does the cognitive load task being learned by the best model also map to social evaluative stress from [**schmidt2018introducing**] or the cognitive load in [**gjoreski2020datasets**]?

Conclusions

This dataset provides a large sampling of subjects with a rigid and well-defined protocol that induces a commonly investigated mental state of cognitive load. This contrasts with common physiological datasets, where the main goal is to induce heightened states of emotion through Trier Stress Tests [**schmidt2018introducing**] or using provocative video [**sharma2019dataset**]. The other issue is that most datasets are of a small sampling of people, usually less than 40 individuals and more commonly around 15. This leads to methodological issues where the generalization performance of the models cannot be assessed in a fair way. The most common way of assessing generalization performance is the Leave One Subject Out approach, which can introduce bias by selecting the best performing subject to report. By having a very large dataset like cLACIr with 140 individuals, even a 10% train-test split leaves 14 individuals in the test set, which is nearly the same size as WESAD [**schmidt2018introducing**].

In addition, this dataset performs well on a variety of machine learning models, showing that the class labels are aligning well with the physiological data features and that learning is occurring. This is further enhanced by using multiple rich metrics that give insight into other aspects of the model's performance.

Ethical Statement

All data was collected and processed in an ethical manner, in full compliance with the UNL Internal Review Board's (protocol #19552) and Institutional Animal Care and Use Committee's (protocol #1599) relevant codes of experimentation and legislation. All participants gave written consent to participate and acknowledged that de-identified data could be published publicly.

Acknowledgments

Author contributions statement

Competing interests